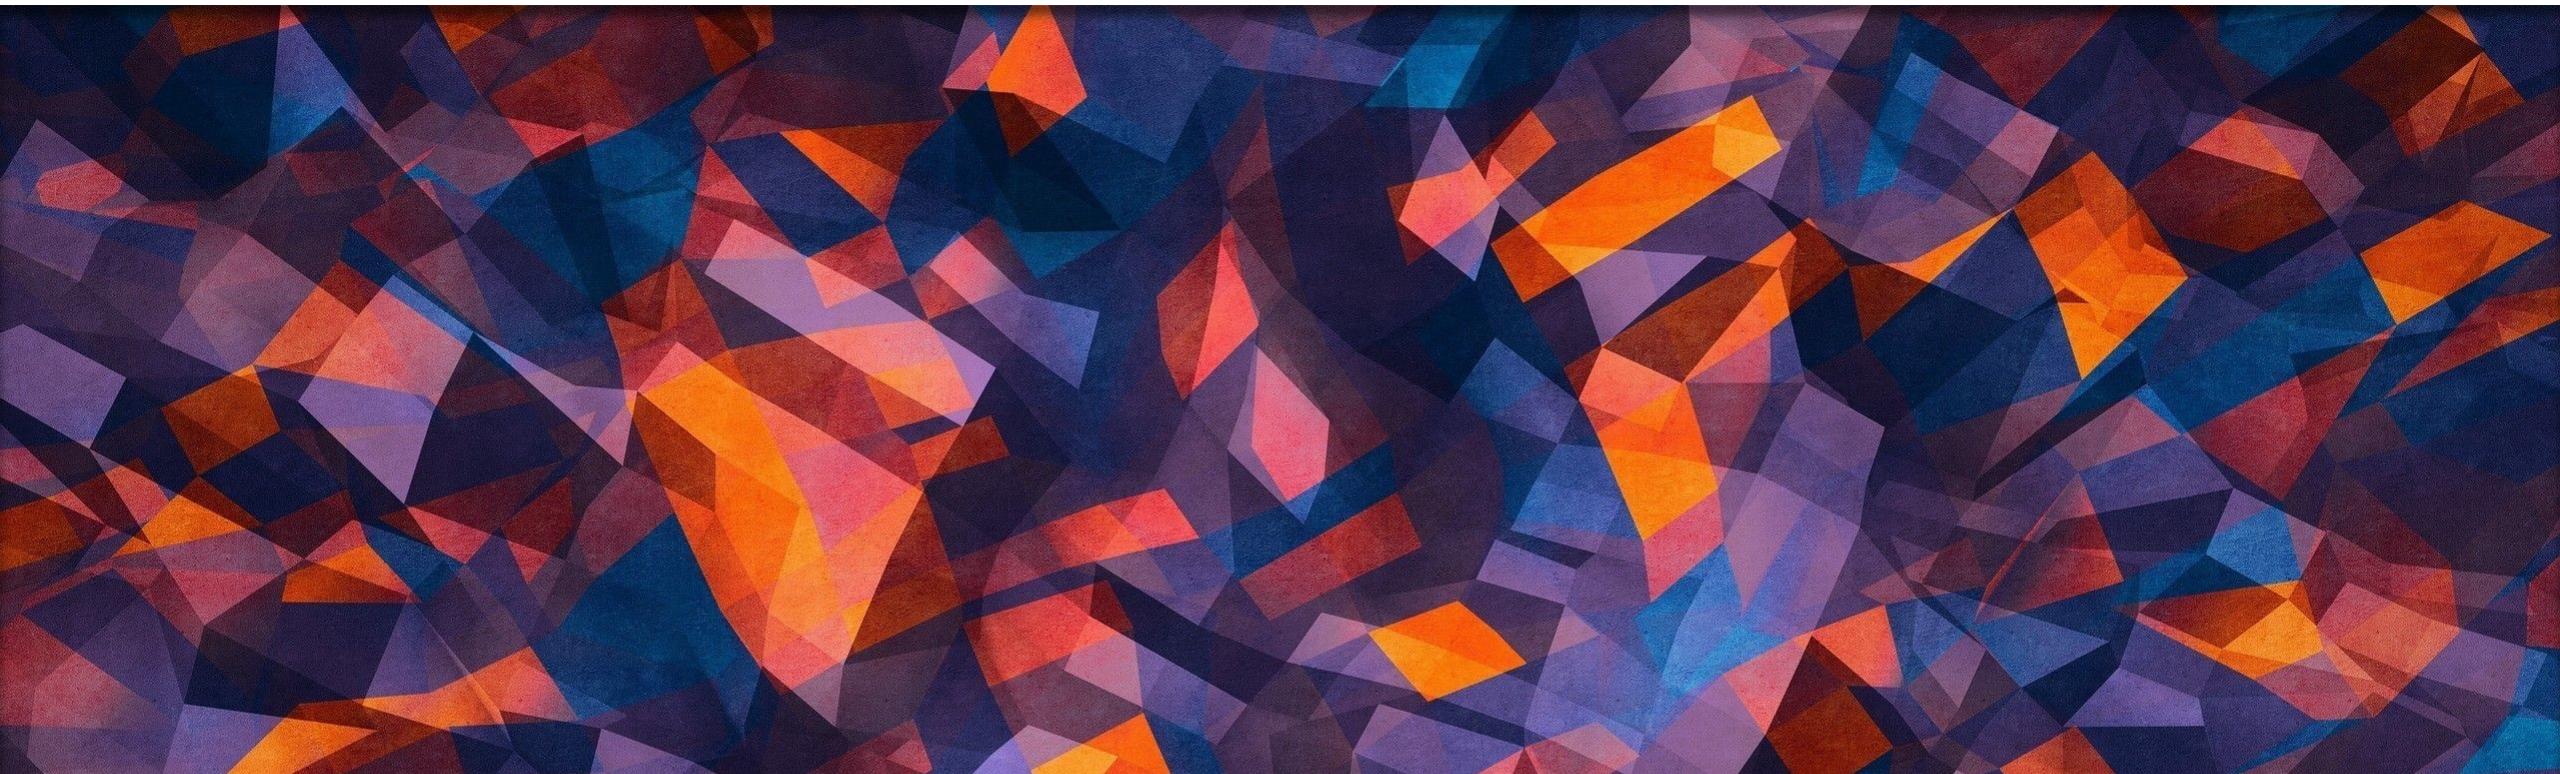


Reflecting the best of human society & evolution (ants)
onto digital systems → rational systems



Web-based Autonomous Systems

Coordination III: Normative Multi-Agent Systems and Organizations

Chair for Interaction- and Communication-based Systems (ICS-HSG)

Our Journey

Prerequisites:
• ASSE
• (...)



Week 1:
Introduction

Week 2:
A Web for Machines

Week 3:
**Knowledge Representation
and Reasoning for the Web**

Week 4:
**Linked Data and Distributed
Knowledge Graphs**

Week 6 (Coordination I):
**Agent Communication
and Interaction**

Week 5:
**Autonomous Agents
(Arch. and Programming)**

Week 7 (Coordination II):
**Self-Organization and
Stigmergy**

Week 9 (Coordination IV):
Trust & Reputation

Week 10:
**Game Theory and
Social Choice**

Week 11:
**Reinforcement Learning
and Multi-Agent Learning**

Week 12:
An Industry Perspective

Exercises

Ex1: Writing Your First Agent(s)!
Ex2: Automated Planning
Ex3: Web Ontologies
Ex4: Operating on Linked Data
Ex5: BDI Agents
Ex6: Interacting Agents on the Web

Ex7: Ant Colony Optimization
Ex8: Organized Agents
Ex9: Trustworthy Agents
Ex10: Axelrod's Agents
Ex11: Reinforcement Learning Agents
Course Review and Q&A

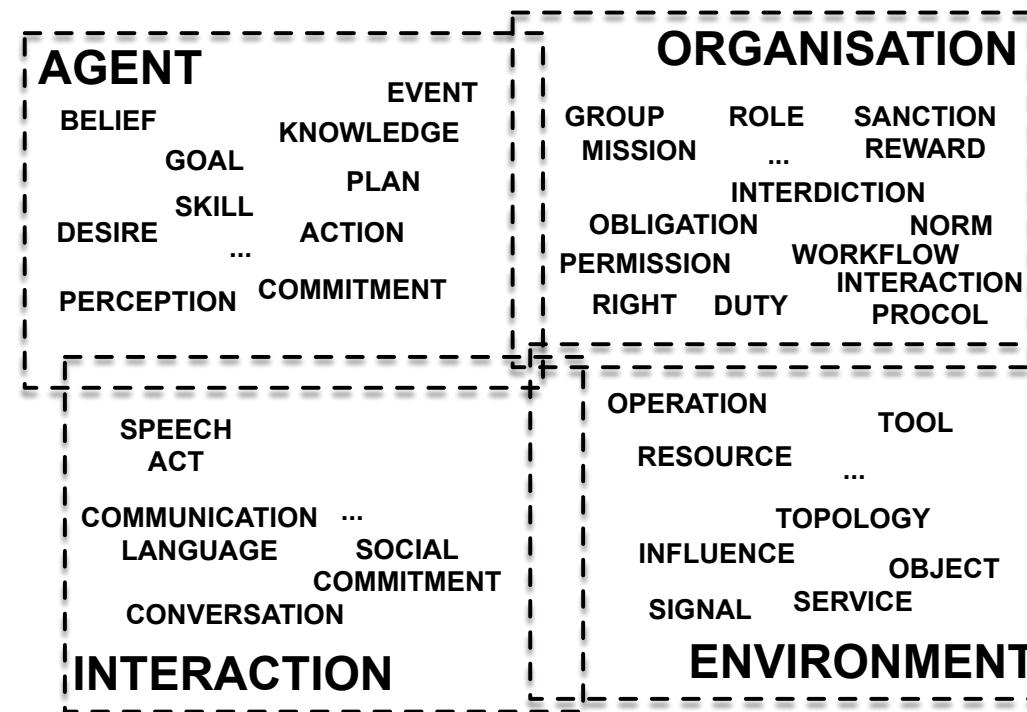
Coordination Perspectives

Week 9 (Coordination IV):
Trust & Reputation and
Multi-Agent Planning

Week 6 (Coordination I):
Communication and
Interaction

Week 8 (Coordination III):
Normative MAS and
Organizations

Week 7 (Coordination II):
Self-Organization and
Stigmergy



Modelling Dimensions for Multi-Agent Systems
[Demazeau, 1995]

O. Boissier, *Multi-Agent Coordination Course*, MINES Saint-Étienne, 2021.

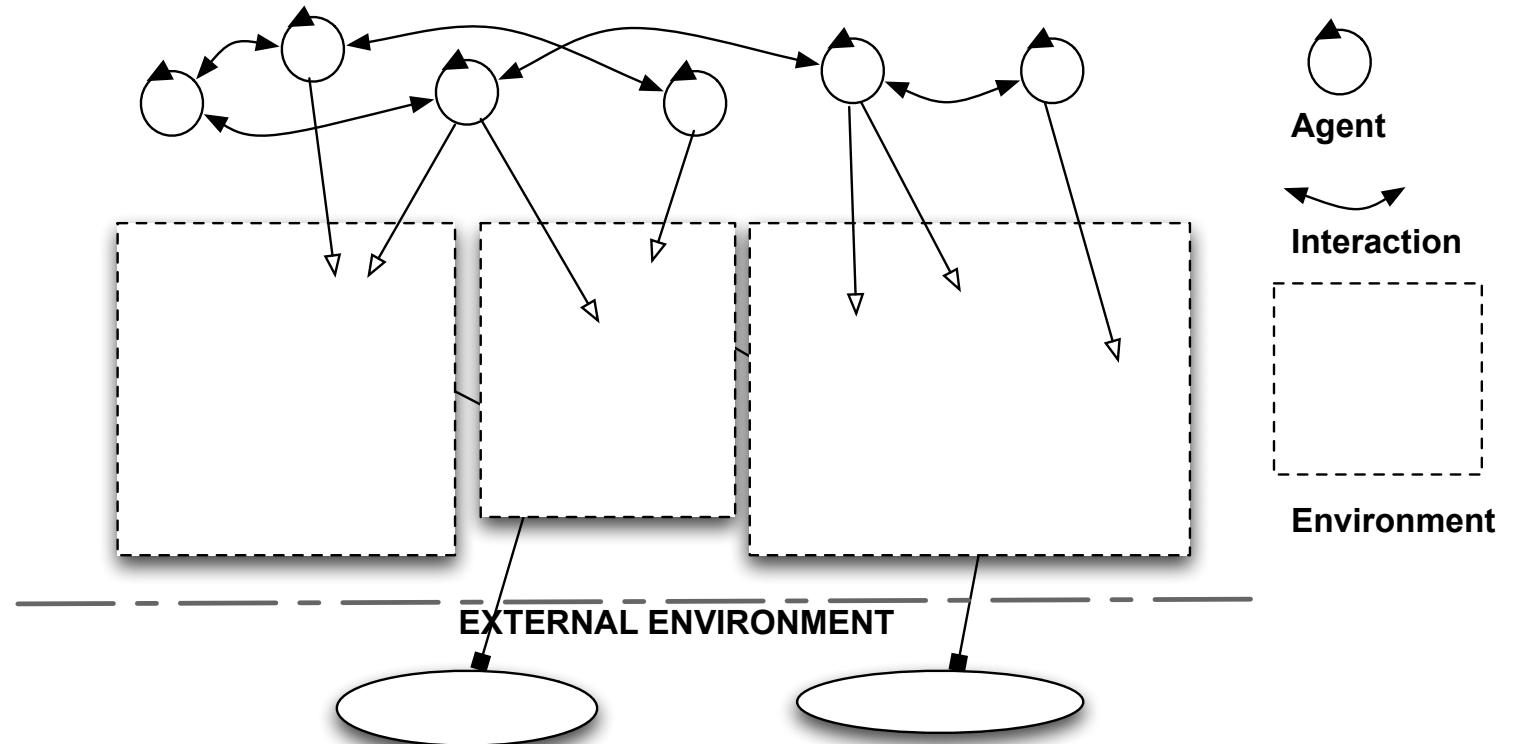
Y. Demazeau. From Interactions To Collective Behaviour In Agent-Based Systems, 1st European Conf. on Cognitive Science, 1995.

Autonomy & Control: Bottom-up vs. Top-down

Local interactions create global emergent behavior

How to enable **control** over the emergent behavior?

How to achieve a desired global behavior **efficiently**?

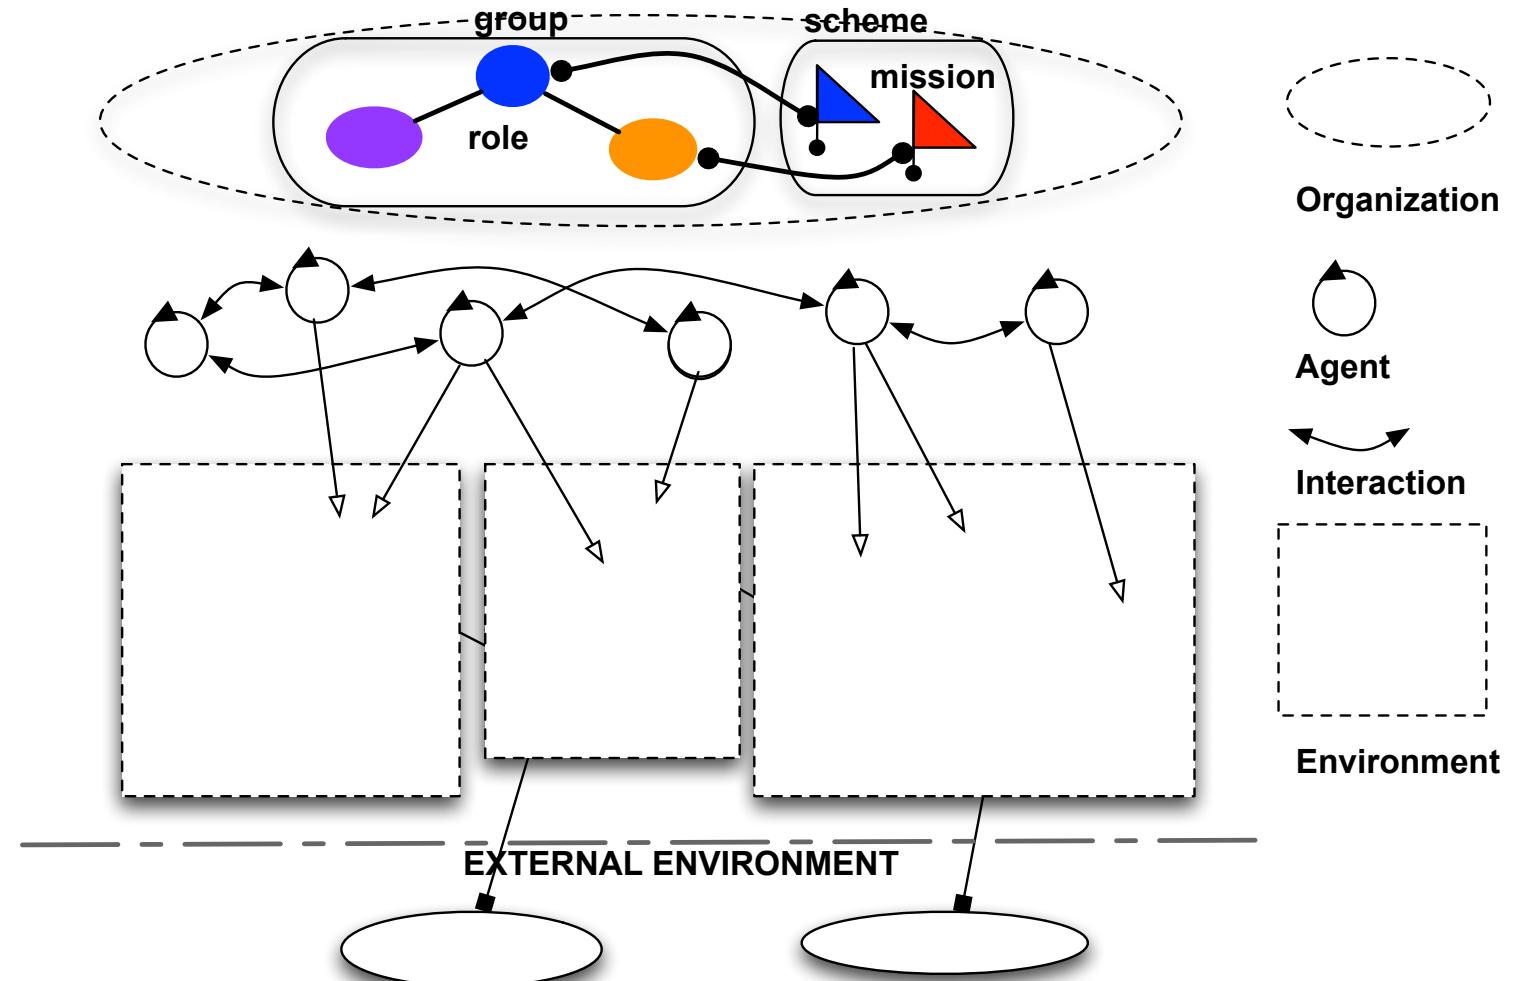


Autonomy & Control: Bottom-up vs. Top-down

Local interactions create global emergent behavior

How to enable **control** over the emergent behavior?

How to achieve a desired global behavior **efficiently**?

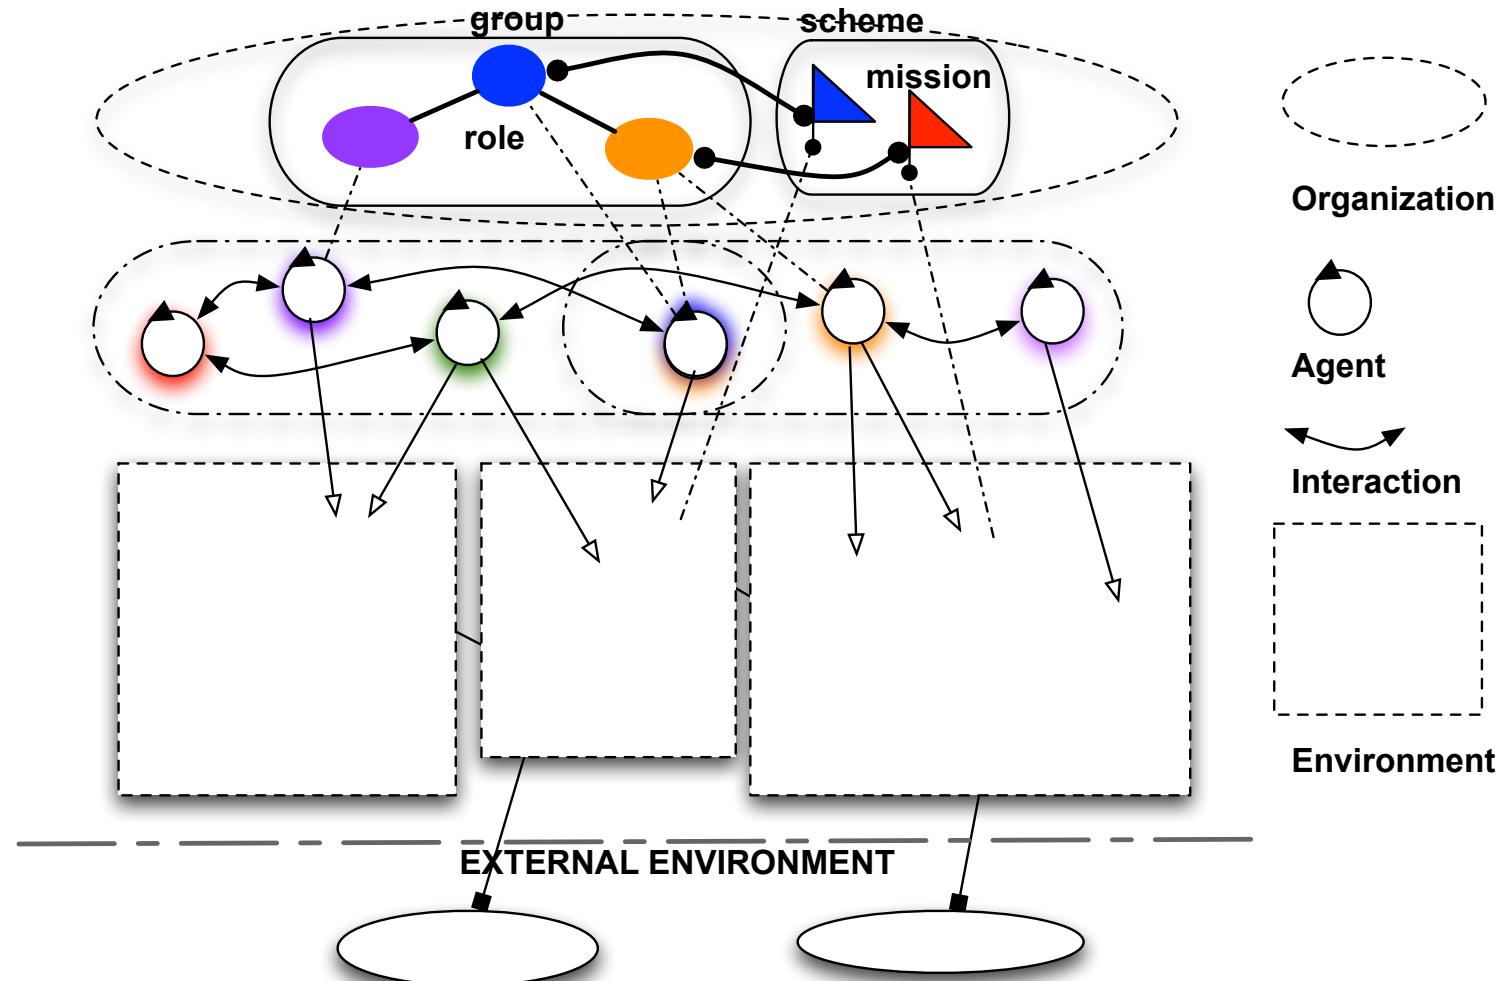


Autonomy & Control: Bottom-up vs. Top-down

Local interactions create global emergent behavior

How to enable **control** over the emergent behavior?

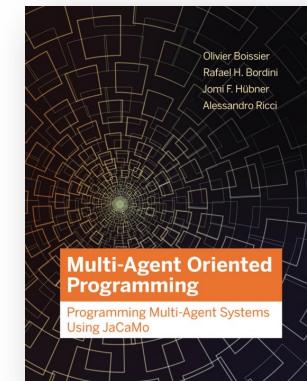
How to achieve a desired global behavior **efficiently**?



Today's Agenda

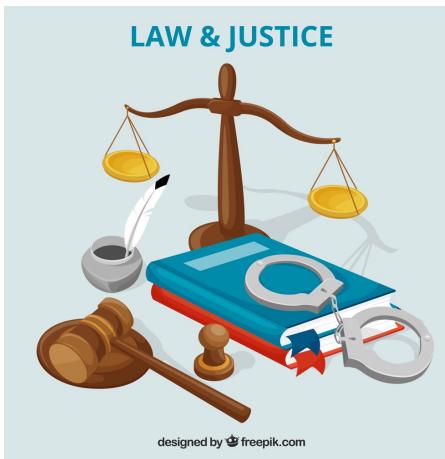
- Normative Multi-Agent Systems
- Normative Organizations
 - Organization-Oriented Programming
 - The Moise Organization Model
 - Hands-on: Programming Organizations with JaCaMo

⇒ revisiting Autonomy from w01/w02



Chapters 8-9

Norms in Human Society



Regulative Norms



Constitutive Norms

raising hand in \sqsubset auction \rightarrow buy
street \rightarrow taxi

A baseline definition [Verhagen et al., 2018]

“Norms are a regularity of intentional behavior within a certain population.”
(i.e., norms provide the framework within which members of a society interact)

Interaction: norms are at the **interface** between the **individual** and **society**

Agency vs. Structure: the **intentional behavior** is a **regular pattern** within a group



Norms in Everyday
Interactions



Normative Multi-Agent Systems

Norms

Norms define a standard of behavior within a society of agents

Norms are **rules** defined within the scope of a society in order to influence the behavior of human and/or artificial agents.

Normative Multi-Agent System (MAS) [Boella et al., 2008]

A MAS organized by means of mechanisms to **represent, communicate, distribute, detect, create, modify, and enforce norms**, and mechanisms to **deliberate about norms and detect norm violation and fulfillment**.

Deductive : Reason from rules (?)

Induction : Can have false positives (!) ~ problems w/ Chat GPT

"Norm" is defined
by OpenAI

O. Boissier, Multi-Agent Programming Course, MINES Saint-Étienne, 2021.

G. Boella, L. Torre, H. and Verhagen. Introduction to the special issue on normative multiagent systems. Autonomous Agents and Multi-Agent Systems, 17(1):1–10. 2008.

Perspectives on Organization in MAS

Organization: purposeful **supra-agent pattern** of pre-defined or emergent cooperation developed by a designer or by the agents themselves [Boissier et al., 2020].

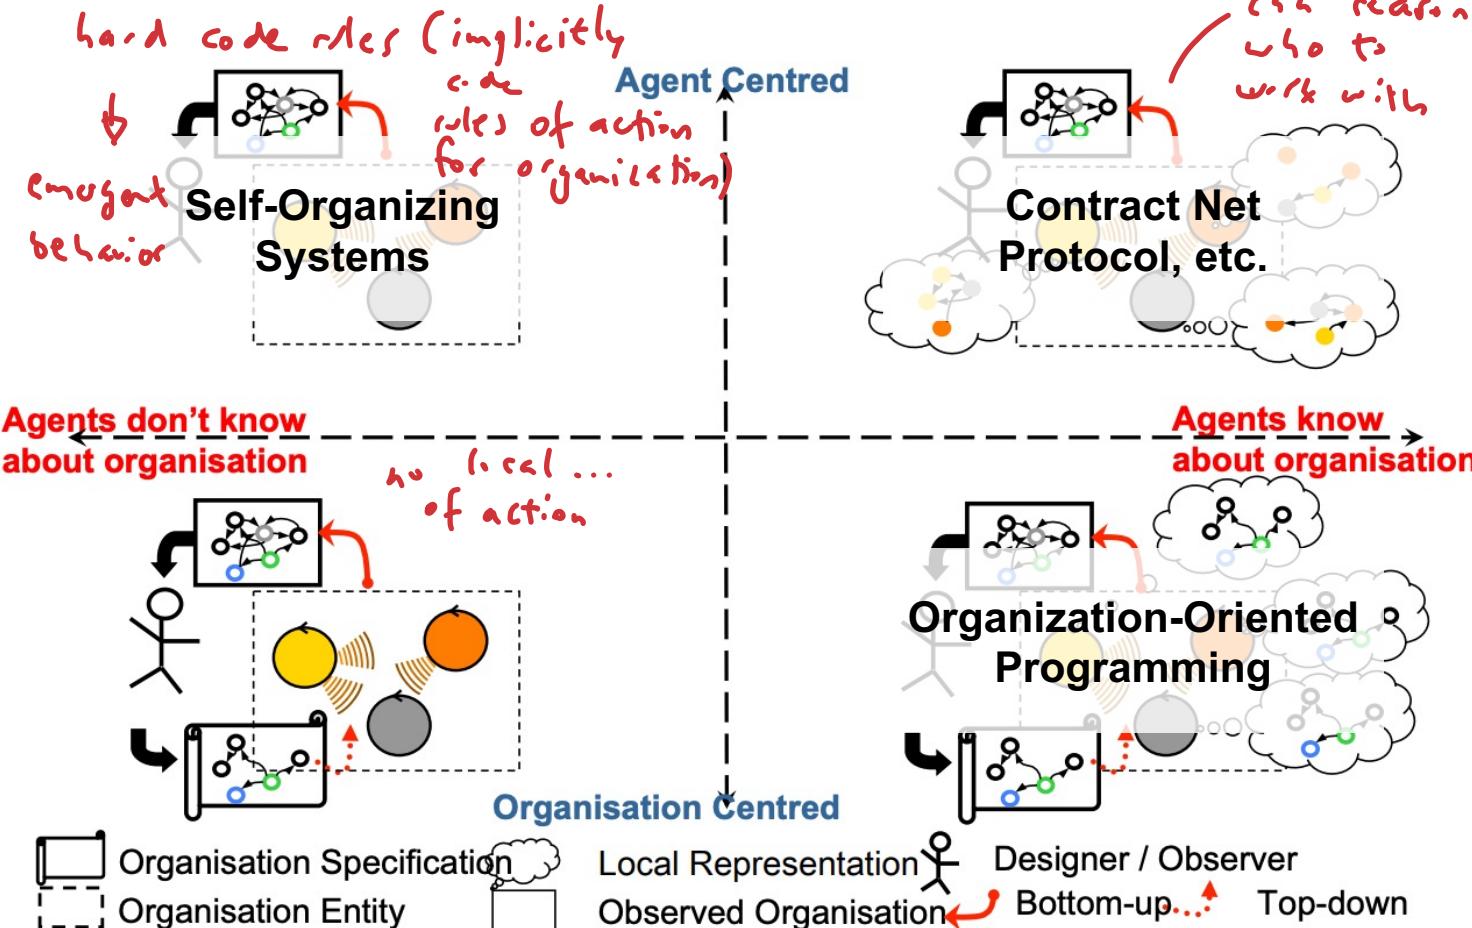
does not have complete picture of end goal

Organization is observed

Organization is programmed implicitly

Organization is a design abstraction

The org. model is hard-coded



Organization is observed

Agents reason to coordinate

Organization is a prog. abstraction

Agents reason on the org. model

→ e.g. what role to adopt

O. Boissier, *Multi-Agent Programming Course*, MINES Saint-Étienne, 2021.

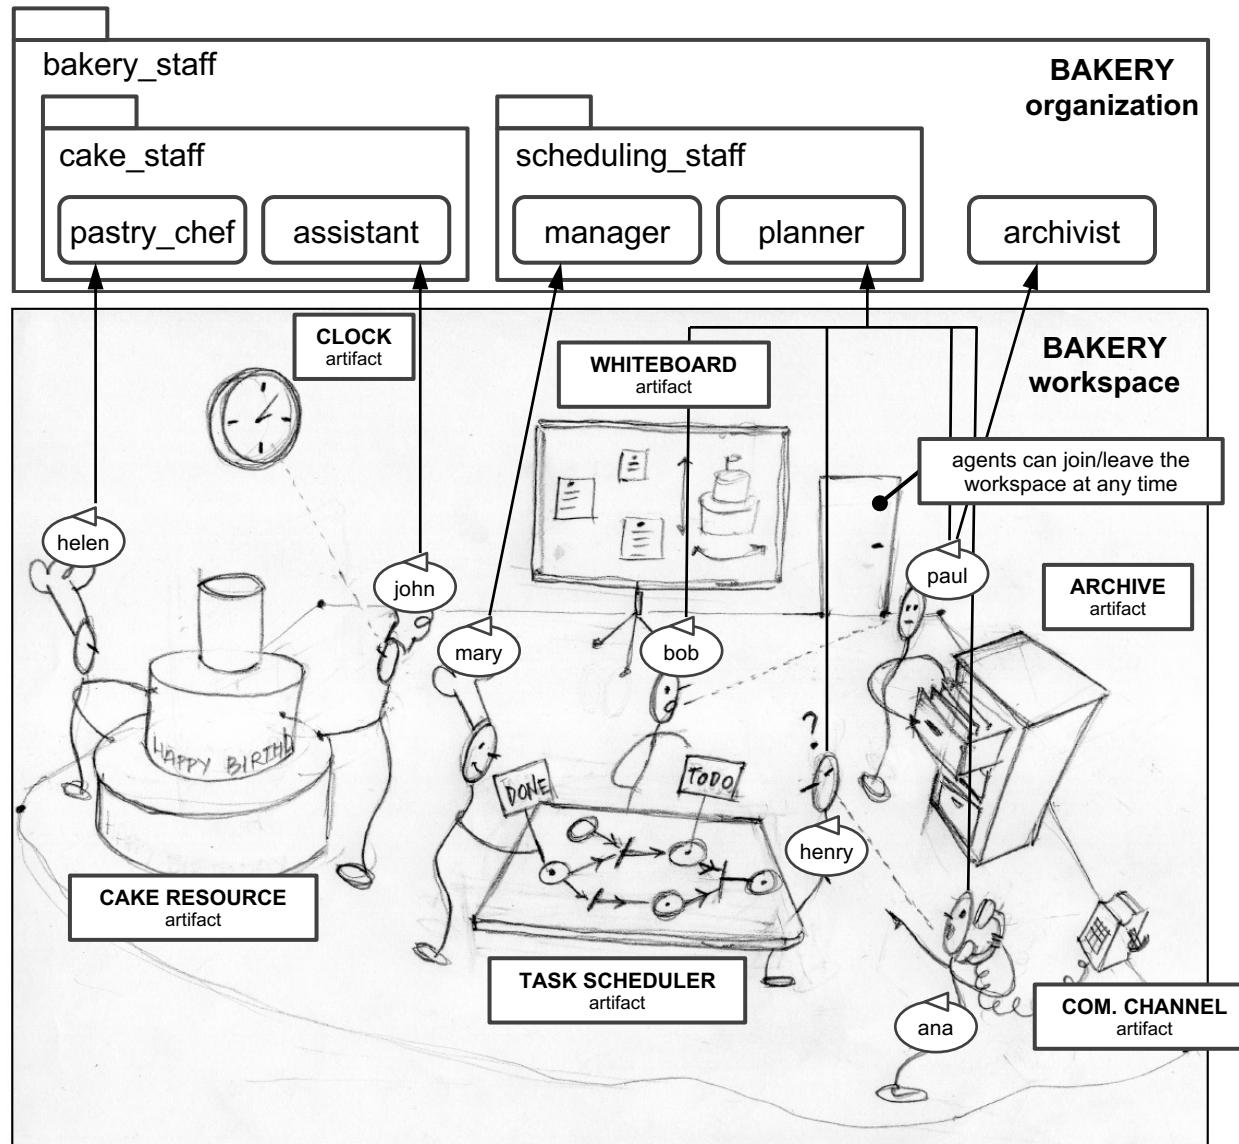
O. Boissier, R. H. Bordini, J.F. Hubner, A. Ricci. *Multi-Agent Oriented Programming: Programming Multi-Agent Systems Using JaCaMo*, The MIT Press, 2020.

Normative Organizations: A First Intuition

decouples role of org from role of agents (via indirection)

In normative organizations:

- norms are used to coordinate agents towards organizational goals
 - ex: to make a birthday cake
- norms are expressed as part of an organization specification
 - ex: “when making a birthday cake, the **assistant** is **obliged** to prepare the workspace”
- norms are interpreted and considered within the context of an organization entity
 - ex: “to make **this birthday cake** for our client, **John** is now **obliged** to prepare the workspace”



How do **normative organizations** impact the **autonomy** of agents?

Review: Autonomy as a Relational Notion

Autonomy as a Relational Notion

A multifaceted view drawing from cognitive science [Castelfranchi, 1993; Castelfranchi and Falcone, 2003]

An entity X is autonomous from Y about P , where:

- X : the main entity whose autonomy is considered/evaluated
- P : a function/action/goal that must be realized or maintained by the main entity and on which the autonomy is evaluated
- Y : the secondary entity (human, artificial agent, environment, organization, etc.) with respect to whom X should be considered autonomous given the specified function/action/goal P

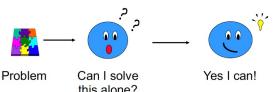
Concrete examples:

- autonomy **from other agents**
- autonomy **from organizations**
- autonomy **from the environment**

need interdependence for an interactive system

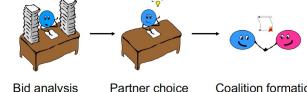
Autonomy from Other Agents

Autonomy as independence



"X is autonomous from (any) Y to achieve goal G"

Autonomy in collaboration



"X depends on Y to achieve goal G"
(and knows Y can help achieve goal G)

Independence vs. Interdependence

- if agents are fully autonomous from other agents, the MAS is just a set of agents that happen to coexist in a shared environment
- autonomous agents must be able to cooperate, collaborate, compete, and carry complex interactions

O. Bösl, Multi-Agent Coordination Course, MFGS Saint-Etienne, 2001
A. Bösl, Multi-Agent Coordination Course, MFGS Saint-Etienne, 2001
R. Falcone, From Automaticity to Autonomy: The Frontier of Artificial Agents. In: Hexmoor H., Castelfranchi C., Falcone R. (eds) Agent Autonomy. Multagent Systems, Artificial Societies, and Simulated Organizations (International Book Series), vol 7. Springer, Berlin, Heidelberg (2003)

- the degree to which an agent's behavior is determined or constrained by the environment

"X is autonomous from the environment to reach destination D"

Applies to virtual environments as well!



Industrial robot on a robot-transfer unit (RTU)

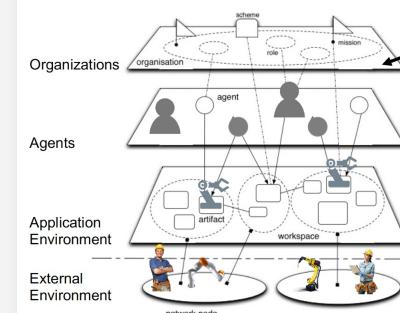


Line-follower robot

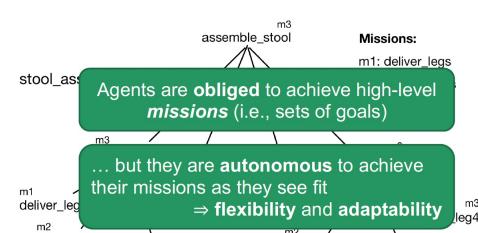
How much is the agent coupled to the environment?

Autonomy from Organizations

"X is permitted/prohibited/obliged by organization Y to achieve goal G"



collectives of people and artificial agents working towards shared goals



Agents are **obliged** to achieve high-level **missions** (i.e., sets of goals)
... but they are **autonomous** to achieve their missions as they see fit
⇒ **flexibility** and **adaptability**

Andrei Ciortea, Simon Mayer, and Florian Michahelles. Repurposing Manufacturing Lines on the Fly with Multi-Agent Systems for the Web of Things, AAMAS 2018.

11

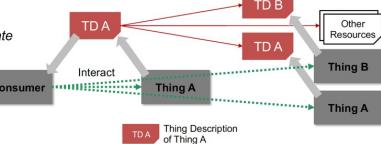
Autonomy from the Environment

Eg TD: Decoupling from environment -> autonomy

Autonomy from the Environment

- the degree to which an agent's behavior is determined or constrained by the environment

Ideally:
arrive-and-operate



TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

Other Resources

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

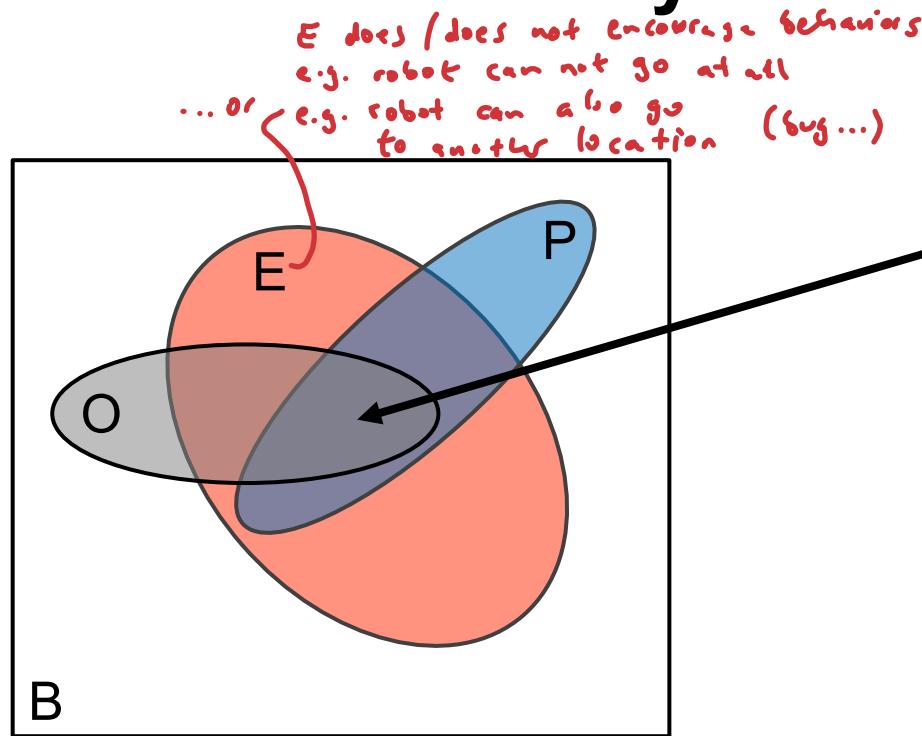
TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

TD B | Thing Description of Thing B

TD A | Thing Description of Thing A

Autonomy and Organizations: A First Intuition



Desirable outcome:

- $P \cap E \cap O$ not too big
 - **constraints** the agents' autonomy
 - increases **performance**
- $P \cap E \cap O$ not too small
 - **preserves** the agents' autonomy
 - increases **flexibility/adaptability**

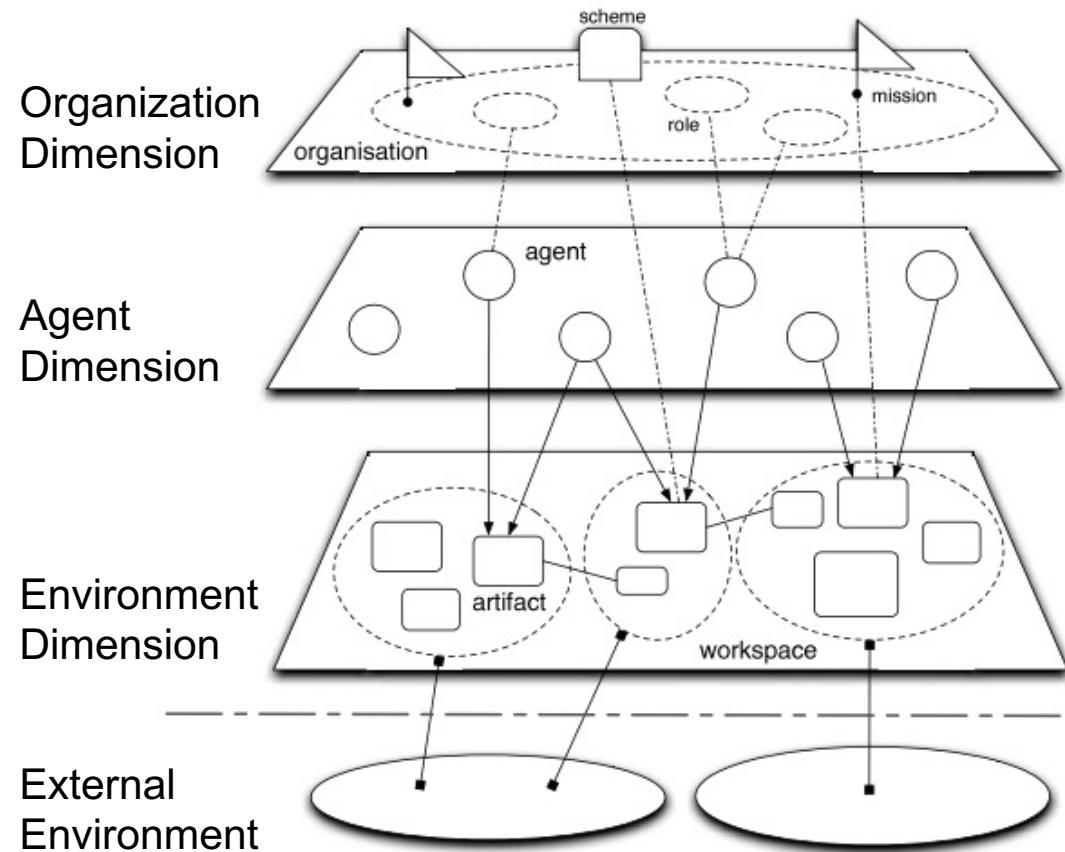
leave enough room to achieve design objective with autonomy

- B: agents' possible behaviors in the MAS
- P: agents' behaviors that lead to some global design objective(s) of the MAS
- E: agents' possible behaviors given environment constraints
- O: agents' possible behaviors given constraints from a normative organization

Today's Agenda

- Normative Multi-Agent Systems
- Normative Organizations
 - Organization-Oriented Programming
 - The Moise Organization Model
 - Hands-on: Programming Organizations with JaCaMo

Engineering Dimensions for MAS



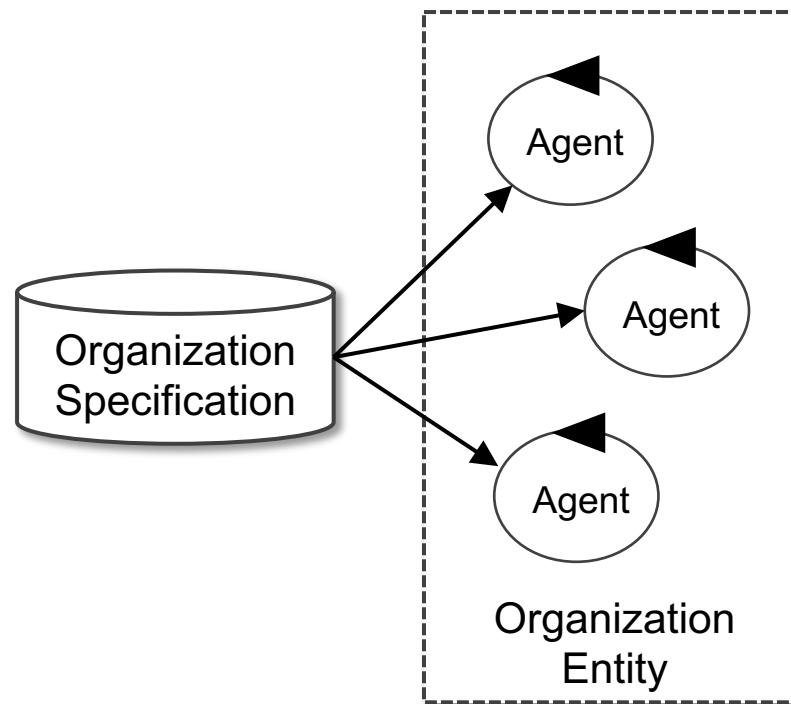
JaCaMo meta-model [Boissier et al., 2013]

Programming MAS = Programming **Agents**
+ Programming the **Environment**
+ Programming the **Organization**

The **environment** is a **first-class design and programming abstraction**

The **organization** is a **first-class design and programming abstraction**

Organization-Oriented Programming



Organization specification:

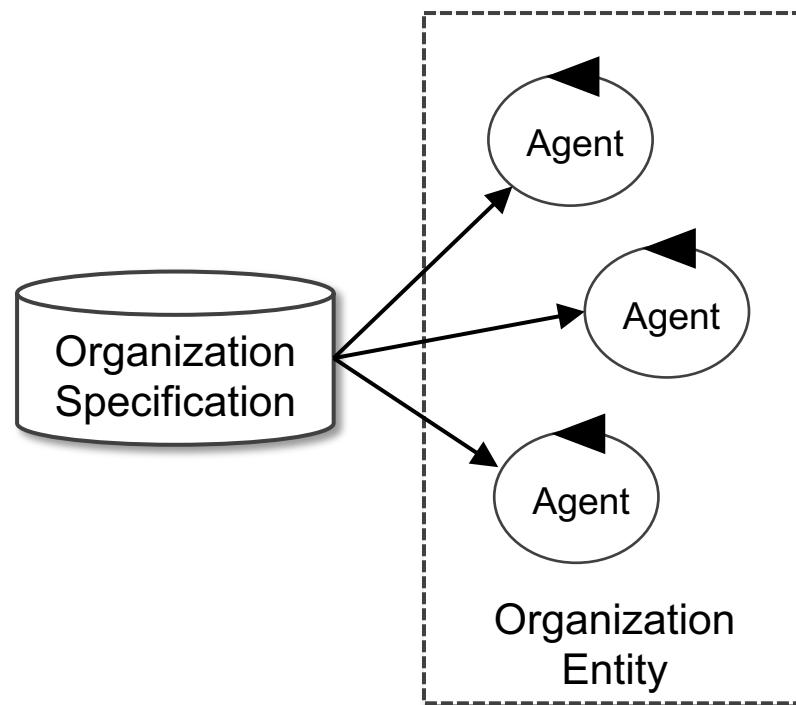
- defines a cooperative pattern using **organizational concepts** such as roles, groups, workflows, norms, etc.
- a **first-class abstraction**: programmed outside of the agents and outside of the environment
- by changing the organization specification, we can change the **overall behavior** of the MAS

Organization entity:

- an enactment of an organization specification
- can be represented in the mind of agents, and/or can be represented explicitly in the MAS

Program =
Organization Specification

Organization-Oriented Programming

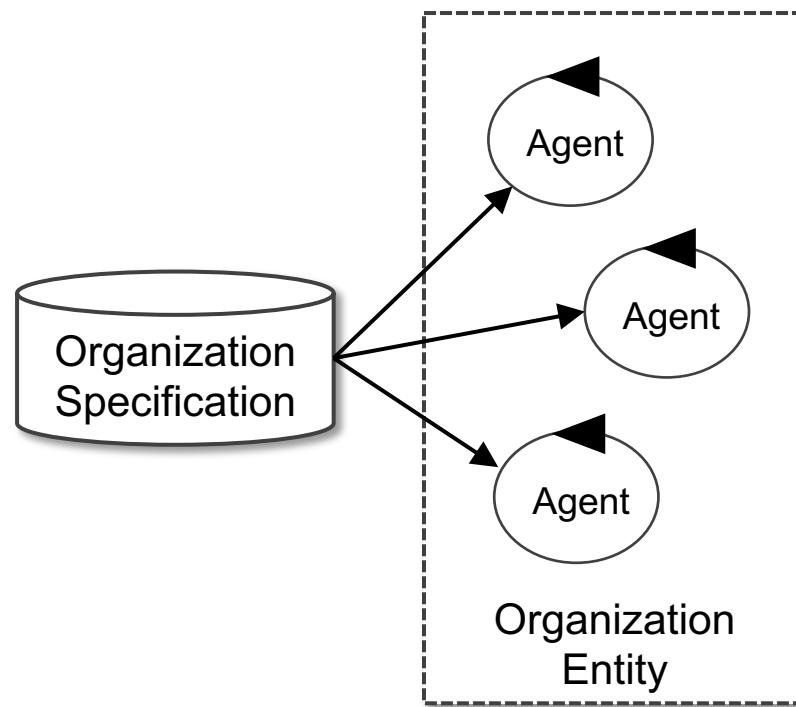


The organization specification is written in a programming language (**organization modelling language**)

Agents **read** the organization specification, **reason** on it, and **follow** the specification (or not)

How to ensure that agents follow an organization specification?

Organization-Oriented Programming

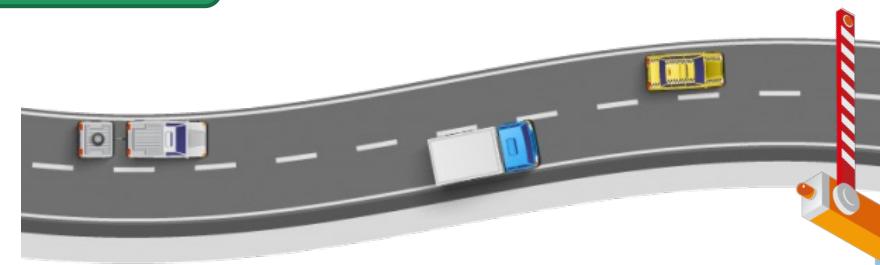


The organization specification is written in a programming language (**organization modelling language**)

Agents **read** the organization specification, **reason** on it, and **follow** the specification (or not)

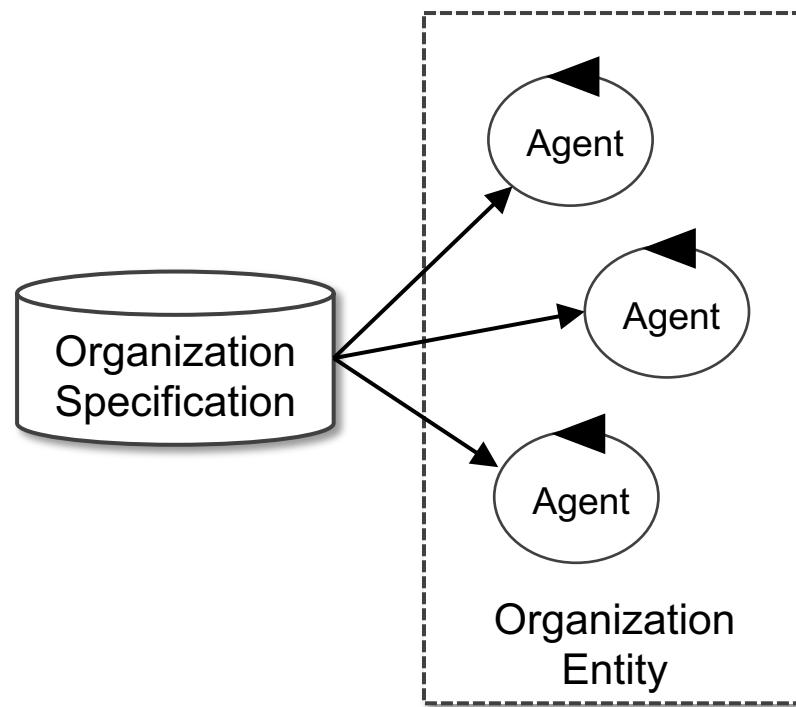
- **regimentation**: agents are **forced** to follow the specification

Lecture #1



Objective: design a system for collecting highway taxes

Organization-Oriented Programming



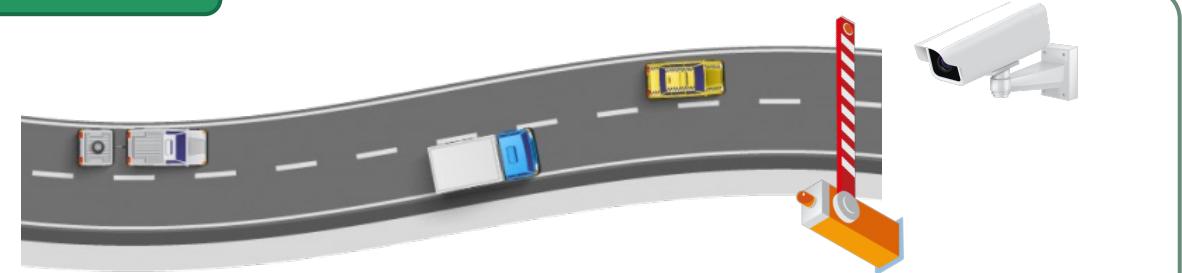
The organization specification is written in a programming language (**organization modelling language**)

Agents **read** the organization specification, **reason** on it, and **follow** the specification (or not)

- **regimentation**: agents are **forced** to follow the specification
- **regulation**: agents are **rewarded** if they follow the specification and/or **sanctioned** if not

Advantages and disadvantages?

Lecture #1

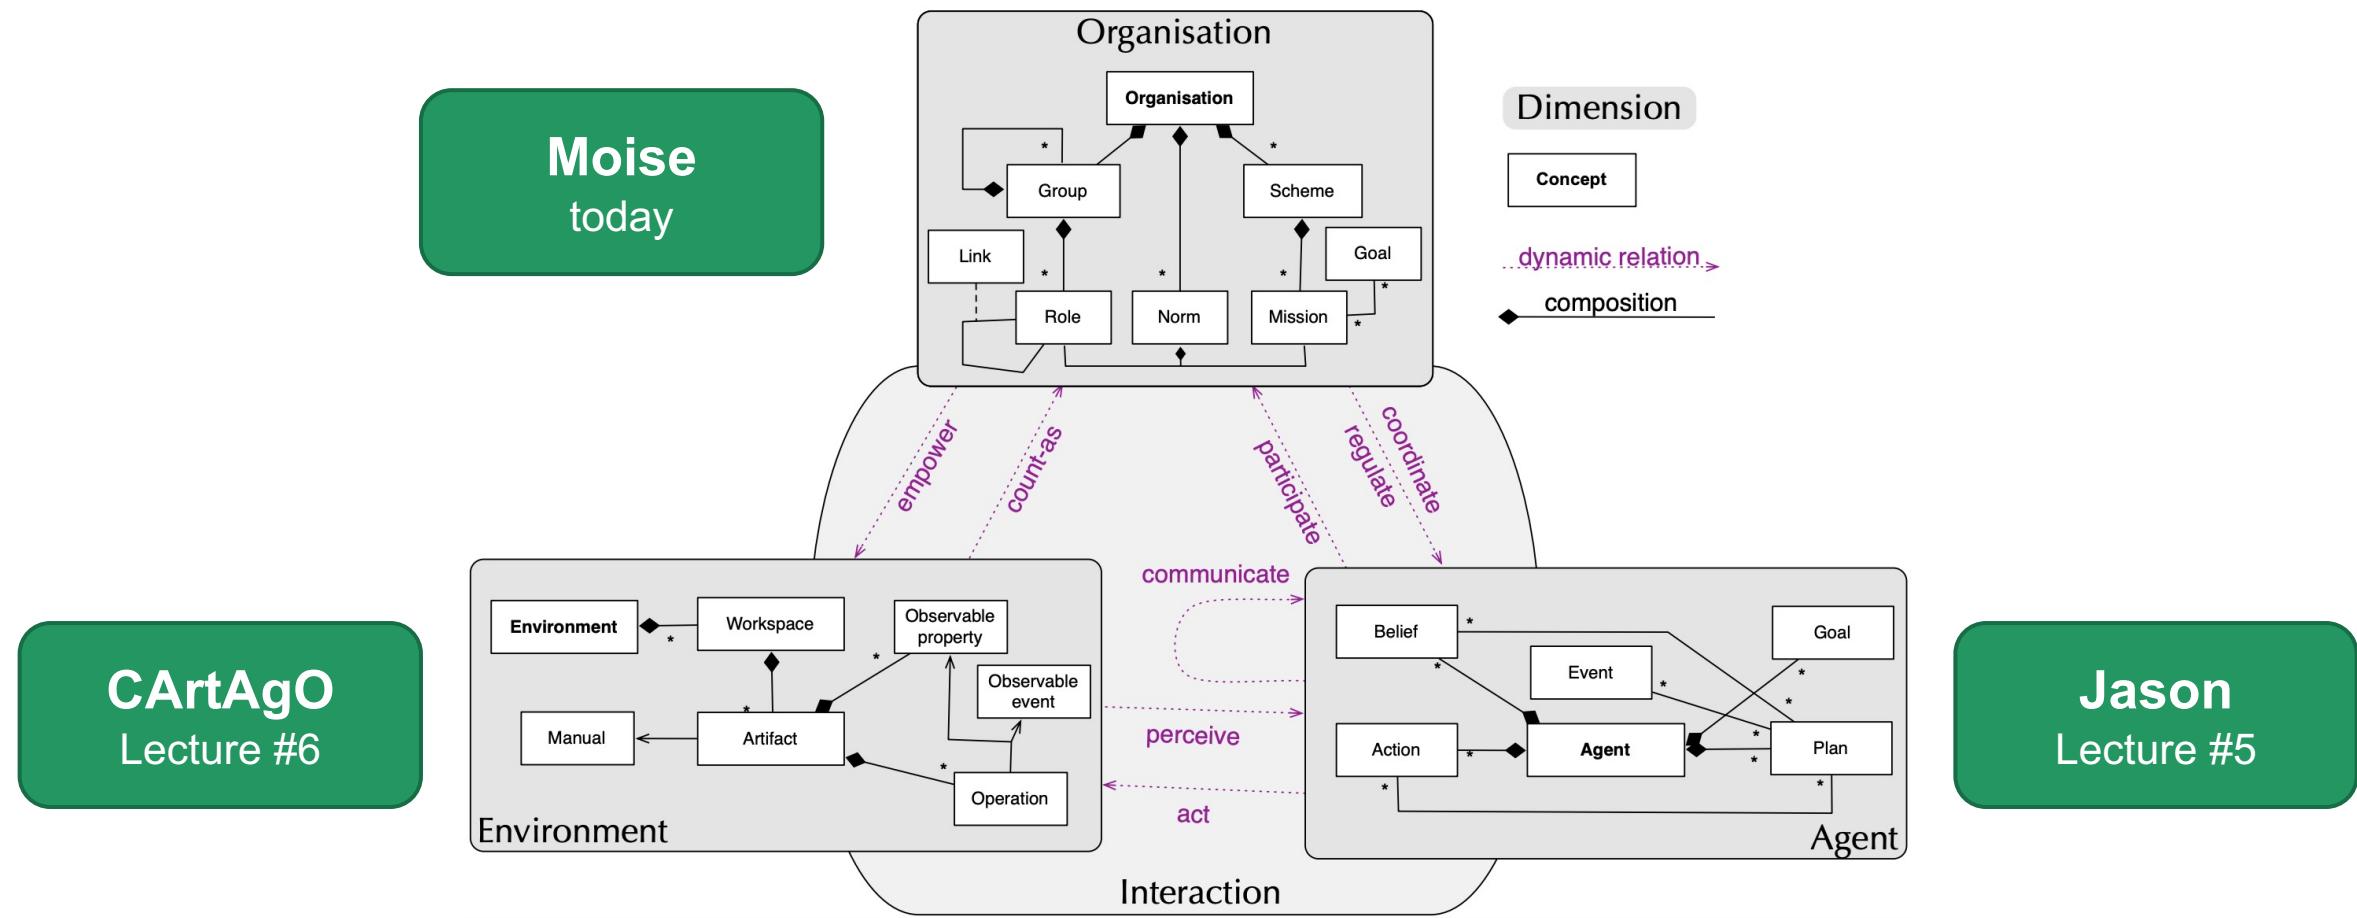


Objective: design a system for collecting highway taxes

Today's Agenda

- Normative Multi-Agent Systems
- Normative Organizations
 - Organization-Oriented Programming
 - The Moise Organization Model
 - Hands-on: Organization-Oriented Programming with JaCaMo

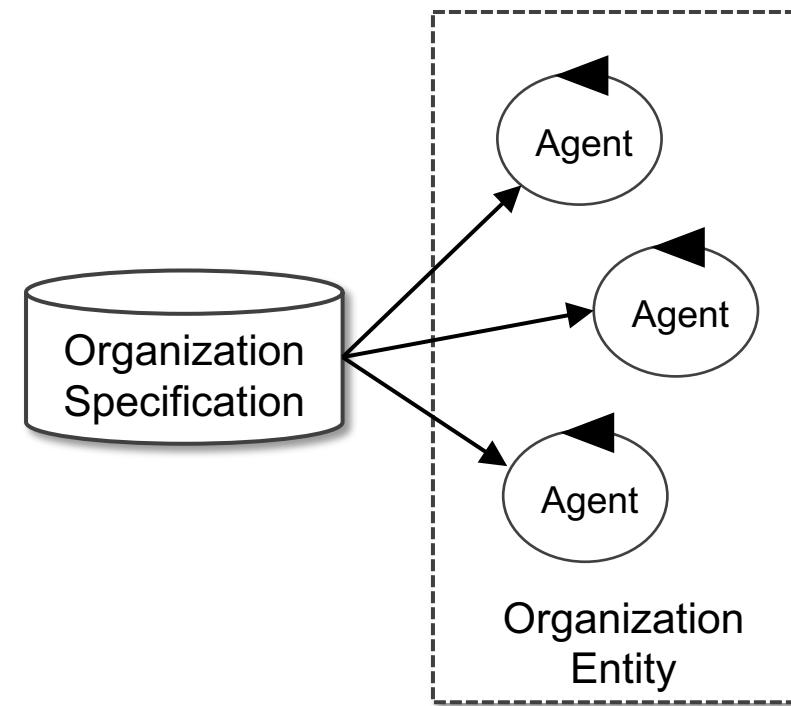
The JaCaMo Meta-Model Revisited



JaCaMo = Jason + CArtAgO + Moise

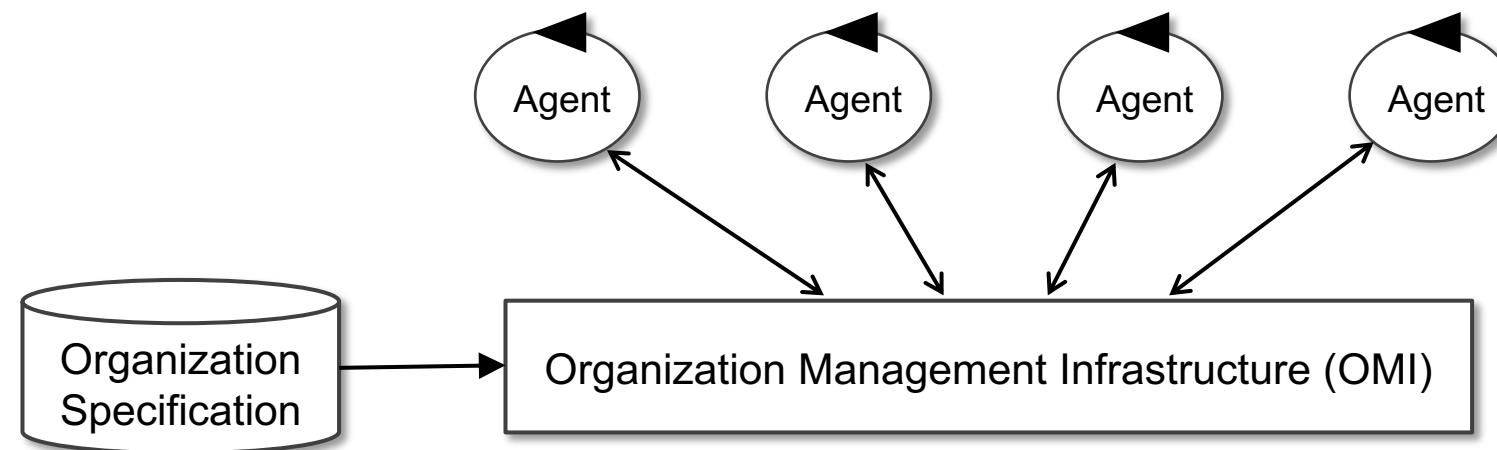
Programming Organizations in JaCaMo

So far, we've looked at **organizations** as first-class programming abstractions



Programming Organizations in JaCaMo

So far, we've looked at **organizations** as first-class programming abstractions
... but organizations are typically managed through a **shared infrastructure**



OMI Responsibility

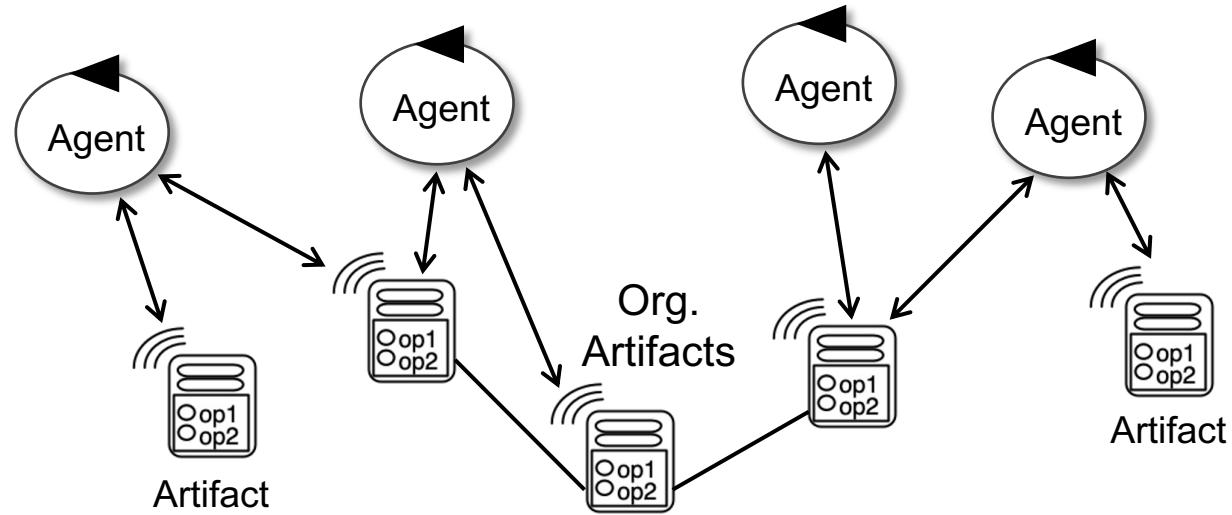
To manage the agents' execution within an organization entity defined by an organization specification

- coordination, regimentation, regulation

Programming Organizations in JaCaMo

So far, we've looked at **organizations** as first-class programming abstractions

... but organizations are typically managed through a **shared infrastructure**



OMI Responsibility

To manage the **agents' execution** within an organization entity defined by an organization specification

- coordination, regimentation, regulation

In JaCaMo, the OMI is distributed across a set of **organizational artifacts**

The agents' working environment is **instrumented** with **organizational artifacts**

- agents **create, handle, and use** org. artifacts (deploy and manage their OMI)
- **perceive** the organization state and violations of norms
- **decide** about organizational actions, sanctions to apply, etc.

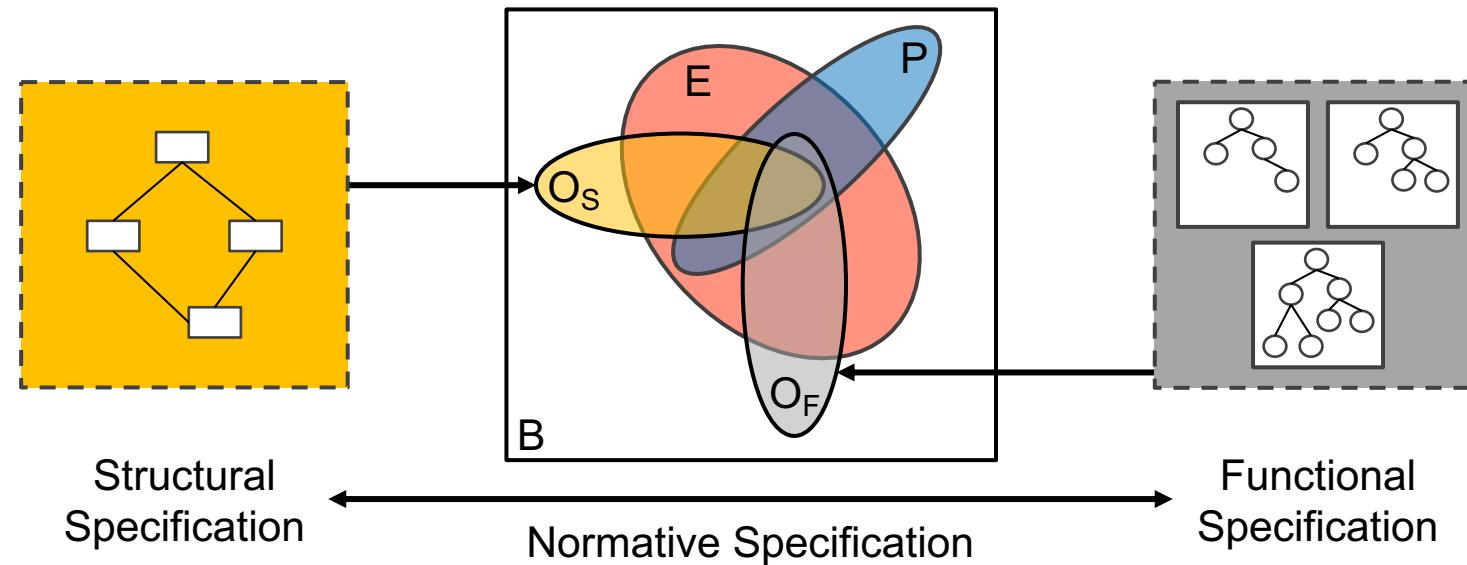
Cognitive Stigmergy

Features supported by the Moise framework

The Moise Framework

Moise defines an **organization specification (OS)** on 3 independent dimensions:

- a **structural specification**: defines structure in terms of **roles** and **groups**
- a **functional specification**: defines workflows in terms of **goals**, **missions**, and **schemes**
- a **normative specification**: flexible coupling between the structural and functional specifications through **norms** that assign missions to roles (**permissions** and **obligations**)



The Moise Framework

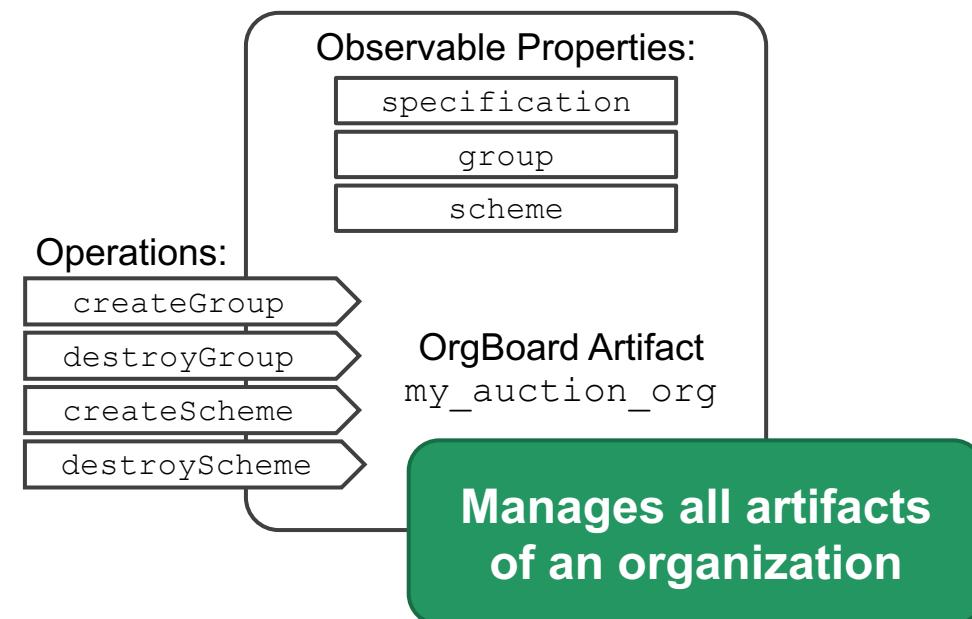
A Moise OS is defined using the **Moise Organization Modelling Language (OML)**, an XML-based language

To create an organization, agents can use the OS to instantiate an `OrgBoard` artifact

```
makeArtifact(my_auction_org, "ora4mas.nopl.OrgBoard", ["src/org/auction-os.xml"], OrgArtId);  
    artifact name           points to an OS      new artifact id (output)
```

OS template in Moise OML

```
<organisational-specification  
  <structural-specification>  
    ...  
  </structural-specification>  
  <functional-specification>  
    ...  
  </functional-specification>  
  <normative-specification>  
    ...  
  </normative-specification>  
</organisational-specification>
```



The Structural Specification (Simplified View)

Main concepts: **Role, Link, Group**

- **Role**: label used to assign constraints on the behavior of agents (an indirection level)
 - a role can **inherit** from one or multiple roles

Example

```
<structural-specification>
  <role-definitions>
    <role id="auctioneer" />
    <role id="participant" />
    <role id="vip_participant">
      <extends role="participant" />
    </role>
  </role-definitions>
  <group-specification id="auction_group">
    <roles>
      <role id="auctioneer" min="1" max="1"/>
      <role id="participant" min="0" max="300"/>
    </roles>
    <links>
      <link from="auctioneer" to="participant"
        type="communication" scope="intra-group"
        extends-subgroups="true" bi-dir="true" />
    </links>
  </group-specification>
</structural-specification>
```

The Structural Specification (Simplified View)

Main concepts: **Role, Link, Group**

- **Role**: label used to assign constraints on the behavior of agents (an indirection level)
 - a role can **inherit** from one or multiple roles
- **Link**: relation between roles that constraints the interaction between agents enacting the roles
 - link types: *acquaintance, communication, authority*

Example

```
<structural-specification>
  <role-definitions>
    <role id="auctioneer" />
    <role id="participant" />
    <role id="vip_participant">
      <extends role="participant" />
    </role>
  </role-definitions>
  <group-specification id="auction_group">
    <roles>
      <role id="auctioneer" min="1" max="1"/>
      <role id="participant" min="0" max="300"/>
    </roles>
    <links>
      <link from="auctioneer" to="participant"
        type="communication" scope="intra-group"
        extends-subgroups="true" bi-dir="true" />
    </links>
  </group-specification>
</structural-specification>
```

The Structural Specification (Simplified View)

Main concepts: **Role**, **Link**, **Group**

- **Role**: label used to assign constraints on the behavior of agents (an indirection level)
 - a role can **inherit** from one or multiple roles
- **Link**: relation between roles that constraints the interaction between agents enacting the roles
 - link types: *acquaintance*, *communication*, *authority*
- **Group**: a set of roles and relations among roles used to define a shared context for agents
 - groups can contain **sub-groups**
 - groups can define **cardinality restrictions** on roles (**min** and **max** attributes)
 - a group is **well-formed** if the cardinality restrictions of all roles are satisfied

Example

```
<structural-specification>
  <role-definitions>
    <role id="auctioneer" />
    <role id="participant" />
    <role id="vip_participant">
      <extends role="participant" />
    </role>
  </role-definitions>
  <group-specification id="auction_group">
    <roles>
      <role id="auctioneer" min="1" max="1"/>
      <role id="participant" min="0" max="300"/>
    </roles>
    <links>
      <link from="auctioneer" to="participant"
        type="communication" scope="intra-group"
        extends-subgroups="true" bi-dir="true" />
    </links>
  </group-specification>
</structural-specification>
```

The GroupBoard Artifact

Manages an instance of a group in the organization entity

- exposes the state of the group via observable properties
- exposes operations for participating in the group

Agents can instantiate a GroupBoard artifact via an OrgBoard artifact using the `createGroup` operation:

```
createGroup(my_auction_group, auction_group, GrpArtId);
```

name of the artifact
to be created

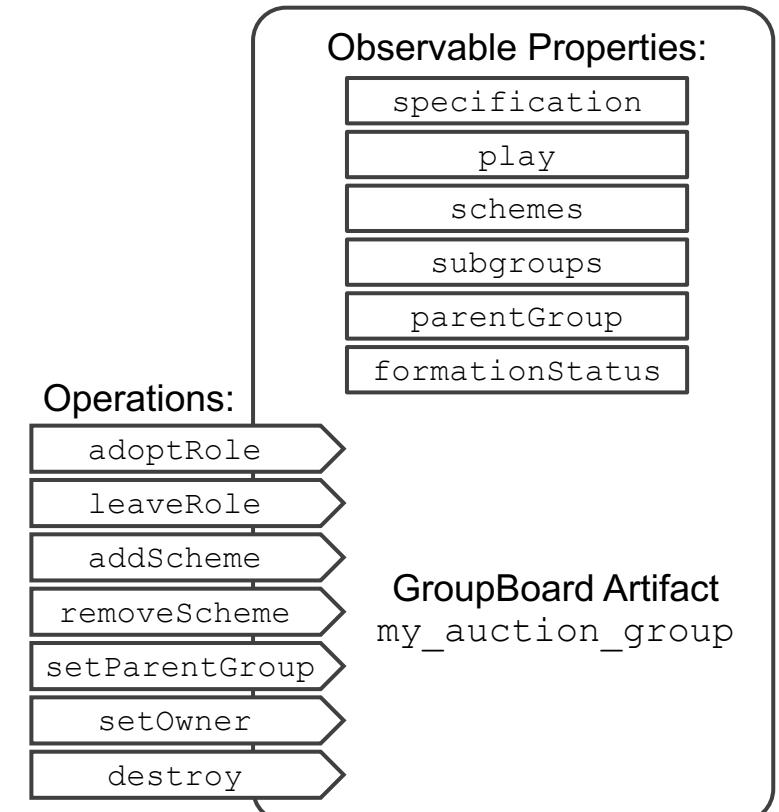
group id
in the OS

id of the artifact
to be created
(output parameter)

or they can instantiate the artifact directly in their workspace:

```
makeArtifact(my_auction_group, "ora4mas.nopl.GroupBoard",  
            ["src/org/auction-os.xml"], GrpArtId);
```

points to a local OS



The Functional Specification (Simplified View)

Main concepts: Goal, Mission, Scheme

- **Goal**: organizational objectives; can be *achievement goals* (default type) or *maintenance goals*

Example

```
<functional-specification>
  <scheme id="do_auction">
    <goal id="auction">
      <argument id="Id" />
      <argument id="Service" />
      <plan operator="sequence">
        <goal id="start_auction" />
        <goal id="bid" ttf="10 seconds" />
        <goal id="decide_winner" ttf="1 hour" />
      </plan>
    </goal>
    <mission id="mAuctioneer" min="1" max="1">
      <goal id="start_auction" />
      <goal id="decide_winner" />
    </mission>
    <mission id="mParticipant" min="1" >
      <goal id="bid" />
    </mission>
  </scheme>
</functional-specification>
```

The Functional Specification (Simplified View)

Main concepts: Goal, Mission, Scheme

- **Goal**: organizational objectives; can be *achievement goals* (default type) or *maintenance goals*
 - goals can have **arguments** (key-value pairs)

Example

```
<functional-specification>
  <scheme id="do_auction">
    <goal id="auction">
      <argument id="Id" />
      <argument id="Service" />
    <plan operator="sequence">
      <goal id="start_auction" />
      <goal id="bid" ttf="10 seconds" />
      <goal id="decide_winner" ttf="1 hour" />
    </plan>
  </goal>
  <mission id="mAuctioneer" min="1" max="1">
    <goal id="start_auction" />
    <goal id="decide_winner" />
  </mission>
  <mission id="mParticipant" min="1" >
    <goal id="bid" />
  </mission>
  </scheme>
</functional-specification>
```

The Functional Specification (Simplified View)

Main concepts: Goal, Mission, Scheme

- **Goal**: organizational objectives; can be *achievement goals* (default type) or *maintenance goals*
 - goals can have **arguments** (key-value pairs)
- **Mission**: a coherent set of goals that can be assigned to a role via a norm

Example

```
<functional-specification>
  <scheme id="do_auction">
    <goal id="auction">
      <argument id="Id" />
      <argument id="Service" />
      <plan operator="sequence">
        <goal id="start_auction" />
        <goal id="bid" ttf="10 seconds" />
        <goal id="decide_winner" ttf="1 hour" />
      </plan>
    </goal>
    <mission id="mAuctioneer" min="1" max="1">
      <goal id="start_auction" />
      <goal id="decide_winner" />
    </mission>
    <mission id="mParticipant" min="1" >
      <goal id="bid" />
    </mission>
  </scheme>
</functional-specification>
```

The Functional Specification (Simplified View)

Main concepts: Goal, Mission, Scheme

- **Goal**: organizational objectives; can be *achievement goals* (default type) or *maintenance goals*
 - goals can have **arguments** (key-value pairs)
- **Mission**: a coherent set of goals that can be assigned to a role via a norm
- **Scheme**: a global goal decomposition tree that can be assigned to a (well-formed) group

Example

```
<functional-specification>
  <scheme id="do_auction">
    <goal id="auction">
      <argument id="Id" />
      <argument id="Service" />
      <plan operator="sequence">
        <goal id="start_auction" />
        <goal id="bid" ttf="10 seconds" />
        <goal id="decide_winner" ttf="1 hour" />
      </plan>
    </goal>
    <mission id="mAuctioneer" min="1" max="1">
      <goal id="start_auction" />
      <goal id="decide_winner" />
    </mission>
    <mission id="mParticipant" min="1" >
      <goal id="bid" />
    </mission>
  </scheme>
</functional-specification>
```

The Functional Specification (Simplified View)

Main concepts: Goal, Mission, Scheme

- **Goal**: organizational objectives; can be *achievement goals* (default type) or *maintenance goals*
 - goals can have **arguments** (key-value pairs)
- **Mission**: a coherent set of goals that can be assigned to a role via a norm
- **Scheme**: a global goal decomposition tree that can be assigned to a (well-formed) group
 - any scheme starts from a **root goal** that is decomposed into sub-goals

Example

```
<functional-specification>
  <scheme id="do_auction">
    <goal id="auction">
      <argument id="Id" />
      <argument id="Service" />
      <plan operator="sequence">
        <goal id="start_auction" />
        <goal id="bid" ttf="10 seconds" />
        <goal id="decide_winner" ttf="1 hour" />
      </plan>
    </goal>
    <mission id="mAuctioneer" min="1" max="1">
      <goal id="start_auction" />
      <goal id="decide_winner" />
    </mission>
    <mission id="mParticipant" min="1" >
      <goal id="bid" />
    </mission>
  </scheme>
</functional-specification>
```

The Functional Specification (Simplified View)

Main concepts: Goal, Mission, Scheme

- **Goal**: organizational objectives; can be *achievement goals* (default type) or *maintenance goals*
 - goals can have **arguments** (key-value pairs)
- **Mission**: a coherent set of goals that can be assigned to a role via a norm
- **Scheme**: a global goal decomposition tree that can be assigned to a (well-formed) group
 - any scheme starts from a **root goal** that is decomposed into sub-goals
 - a non-leaf goal is decomposed through **plans** with one of 3 operators: *sequence*, *choice*, *parallel*

Example

```
<functional-specification>
  <scheme id="do_auction">
    <goal id="auction">
      <argument id="Id" />
      <argument id="Service" />
      <plan operator="sequence">
        <goal id="start_auction" />
        <goal id="bid" ttf="10 seconds" />
        <goal id="decide_winner" ttf="1 hour" />
      </plan>
    </goal>
    <mission id="mAuctioneer" min="1" max="1">
      <goal id="start_auction" />
      <goal id="decide_winner" />
    </mission>
    <mission id="mParticipant" min="1" >
      <goal id="bid" />
    </mission>
  </scheme>
</functional-specification>
```

Main concepts: Goal, Mission, Scheme

- **Goal**: organizational objectives; can be *achievement goals* (default type) or *maintenance goals*
 - goals can have **arguments** (key-value pairs)
- **Mission**: a coherent set of goals that can be assigned to a role via a norm
- **Scheme**: a global goal decomposition tree that can be assigned to a (well-formed) group
 - any scheme starts from a **root goal** that is decomposed into sub-goals
 - a non-leaf goal is decomposed through **plans** with one of 3 operators: *sequence*, *choice*, *parallel*
 - schemes can define **cardinality restrictions** on missions (**min** and **max** attributes)
 - a scheme is **well-formed** if the cardinality restrictions of all missions are satisfied

Example

```
<functional-specification>
  <scheme id="do_auction">
    <goal id="auction">
      <argument id="Id" />
      <argument id="Service" />
      <plan operator="sequence">
        <goal id="start_auction" />
        <goal id="bid" ttf="10 seconds" />
        <goal id="decide_winner" ttf="1 hour" />
      </plan>
    </goal>
    <mission id="mAuctioneer" min="1" max="1">
      <goal id="start_auction" />
      <goal id="decide_winner" />
    </mission>
    <mission id="mParticipant" min="1" >
      <goal id="bid" />
    </mission>
  </scheme>
</functional-specification>
```

The SchemeBoard Artifact

Manages an instance of a scheme in the organization entity

- exposes the state of the scheme via observable properties
- exposes operations for participating in the scheme

Agents can instantiate a SchemeBoard artifact via an OrgBoard artifact using the `createScheme` operation:

```
createScheme(my_auction_scheme, do_auction, SchArtId);
```

name of the artifact
to be created

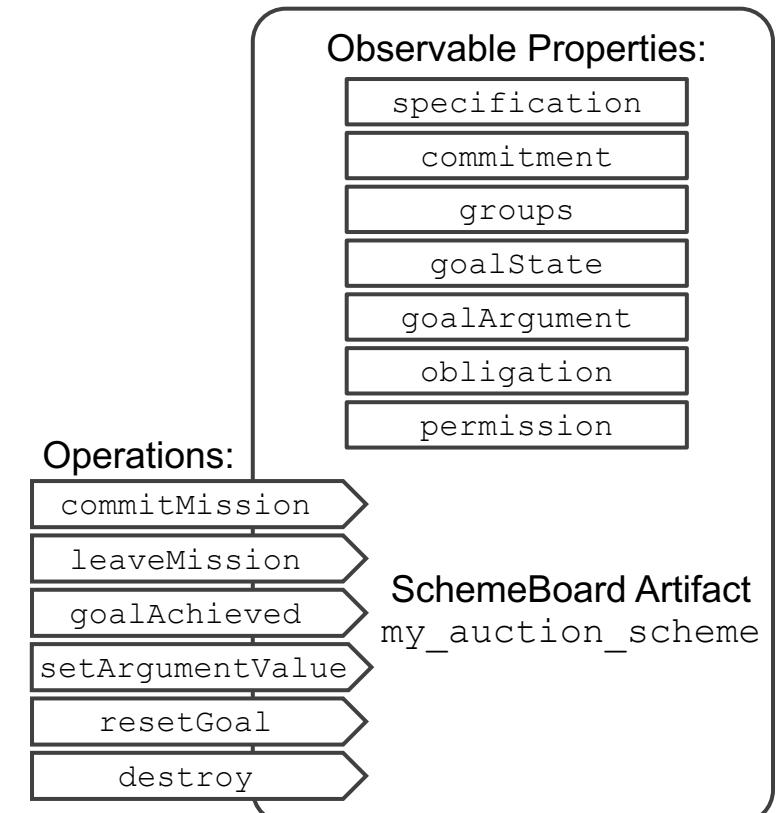
scheme id
in the OS

id of the artifact
to be created
(output parameter)

or they can instantiate the artifact directly in their workspace:

```
makeArtifact(my_auction_scheme, "ora4mas.nopl.SchemeBoard",  
["src/org/auction-os.xml"], SchArtId);
```

points to a local OS



The Normative Specification (Simplified View)

Defines the relation between the structural and functional specifications (**flexible coupling**)

Example

```
<normative-specification>
    <norm id="n1" type="permission"
        role="auctioneer" mission="mAuctioneer" />
    <norm id="n2" type="obligation"
        role="participant" mission="mParticipant" />
</normative-specification>
```

The Normative Specification (Simplified View)

Defines the relation between the structural and functional specifications (**flexible coupling**)

Main concept:

- **Norm**: a coherent assignment of a mission to a role

Example

```
<normative-specification>
  <norm id="n1" type="permission"
        role="auctioneer" mission="mAuctioneer" />
  <norm id="n2" type="obligation"
        role="participant" mission="mParticipant" />
</normative-specification>
```

The Normative Specification (Simplified View)

Defines the relation between the structural and functional specifications (**flexible coupling**)

Main concept:

- **Norm**: a coherent assignment of a mission to a role
 - norm types: *permission, obligation*
 - any mission-to-role assignment that is not an explicit permission or obligation is a prohibition

The normative specification makes explicit the normative dimension of a role

Example

```
<normative-specification>
  <norm id="n1" type="permission"
    role="auctioneer" mission="mAuctioneer" />
  <norm id="n2" type="obligation"
    role="participant" mission="mParticipant" />
</normative-specification>
```

Reasoning About Roles to Adopt

The Moise framework comes with several useful inference rules for reasoning about an OS:

- `role_mission(R, S, M)`: true when there is a norm in the OS that obliges or permits role `R` to commit to mission `M` in scheme `S`
- `mission_goal(M, G)`: true when goal `G` belongs to mission `M`

We can use these 2 rules to write a new rule that determines if a goal `G` is relevant for a role `R`:

```
role_goal(R, G) :-  
    role_mission(R, _, M) & mission_goal(M, G).
```

Using the internal action `.relevant_plans`, we can then write a new rule that determines if the agent has plans for achieving a goal `G`:

```
can_achieve(G) :-  
    .relevant_plans({+!G[scheme(_)]}, LP) & LP \== [].
```

We can now write a rule that returns true if the agent has plans for all goals relevant to role `R`:

```
i_have_plans_for(R) :-  
    not(role_goal(R, G) & not can_achieve(G)).
```

Our Journey

Prerequisites:
 • ASSE
 • (...)



**Week 1:
Introduction**



**Week 2:
A Web for Machines**



**Week 3:
Knowledge Representation
and Reasoning for the Web**



**Week 4:
Linked Data and Distributed
Knowledge Graphs**



**Week 5:
Autonomous Agents
(Arch. and Programming)**



**Week 6 (Coordination I):
Agent Communication
and Interaction**



**Week 7 (Coordination II):
Self-Organization and
Stigmergy**



**Week 9 (Coordination IV):
Trust & Reputation**



**Week 10:
Game Theory and
Social Choice**



**Week 11:
Reinforcement Learning
and Multi-Agent Learning**



**Week 12:
An Industry Perspective**



Exercises

Ex1: Writing Your First Agent(s)!
 Ex2: Automated Planning
 Ex3: Web Ontologies
 Ex4: Operating on Linked Data
 Ex5: BDI Agents
 Ex6: Interacting Agents on the Web

Ex7: Ant Colony Optimization
 Ex8: Organized Agents
 Ex9: Trustworthy Agents
 Ex10: Axelrod's Agents
 Ex11: Reinforcement Learning Agents
 Course Review and Q&A

Any Questions / Comments / Doubts / Concerns?



<https://www.istockphoto.com/>

<https://freepik.com>